

Overview

The instruction set of the CR16 uses standard 16-bit integer ALU operations. So, if we assume that the decoding has been done (actually you will design the decoder in a following lab), what you need to build to execute each instruction is an ALU to compute the numbers, and a register file to provide arguments and hold results. So, this checkpoint has you design and demonstrate those pieces.

As an aside, it's interesting to note how much of the machine can be designed before knowing the complete picture of what the machine will look like. There are a set of basic things that are common to a number of different designs. Building these basic components (memory, register file, ALU, etc) doesn't really require knowing much more than the word size and the overall plan. Filling in the details later will determine the precise character of the final machine. Most VLSI design systems, for example, have large libraries of these generic parts which can be used to design a wide variety of different machine types.

Note: This checkpoint, and all subsequent checkpoints, refer to the *baseline instruction set* for the class processor. You can see the details of the CS/EE 3710 baseline ISA on the class Canvas page. The instructions marked with * are the required baseline subset. The instructions marked with a *? are very strongly encouraged instructions. Basically your group will have to argue convincingly that your application will never use that instruction to get out of implementing it. The instructions with no marking are optional, and probably most groups won't deal with them. The optional instructions mostly deal with exceptions and system processor state, and I expect that most people can get by with polling in their programs and won't need the extra complexity of interrupts. However, interrupts are a very interesting way of handling real-time code, so you might want to explore interrupts, or at least not rule them out quite yet. The blank spots in the instruction encoding are places where you could add any instruction that you would find useful for your specific application and that is not included here.

Execution Datapath

Looking at the Baseline instruction set you should notice a few things:

- The word size is 16 bits. Therefore, although it's not absolutely essential, it would be natural to have a 16 bit ALU. It's not required because you could clearly do the ALU operations with a smaller ALU in a sequential way. The 16-bit result could, for example, be generated using an 8-bit ALU and two timesteps. Going to the extreme, you could use a single-bit ALU and iterate 16-times for each ALU operation. This is not just a crazy hypothetical - there are actually machines that have been built that use bit-serial arithmetic.

Anyway, getting back to the point, your ALU should probably be 16-bits wide to match the data word size.

- The instruction set supports both ADD and SUB instructions, so your ALU should also support these operations. Moreover, these operations can be performed over *unsigned* as well as *signed* integers depending on which instruction is executing. Reading the CR16 handout carefully you will note that *signed* numbers are represented as *two's complement* numbers. So, your ALU should operate on two's complement, and can use this representation for implementing subtraction. The main difference between a *two's complement* ALU and *unsigned* ALU operation is in the way condition codes are generated. Your ALU needs

to generate condition codes for both assumptions. That is, the ALU doesn't know if the bit patterns it sees on the input are unsigned or two's complement, so it has to generate both signed and unsigned condition codes on every operation and let the programmer sort things out by choosing to look at one code or the other.

The condition codes required by the Baseline instruction set are defined here and also in the ISA document. Note that the ADDU and ADDUI instructions are just like ADD and ADDI but they don't set any condition codes.

C — The Carry bit determines whether a carry or borrow occurred after addition or subtraction if the arguments are considered *unsigned*. 0 means no carry or borrow, 1 means a carry or borrow occurred.

L — The Low flag is set by comparison operations. L is set to 1 if the Rdest operand is less than the Rsrc operand when they are both interpreted as *unsigned* numbers.

F — The Flag bit is used by arithmetic operations to signal arithmetic overflow (this is sometimes called the V bit on other machines). This bit should be set if there's a signed overflow on a *signed* addition or subtraction.

Z — The Z bit is set by the comparison operation. It is set to 1 if the two operands are equal, and is cleared otherwise. Here is one extension: The moment the ALU result is 0 (say, due to AND with 0s), why not just set it? Why restrict it to just comparison operations?

N — The Negative bit is set by the comparison operation. It is set to 1 if the Rdest operand is less than the Rsrc operand when both are considered to be *signed* integers.

- The processor supports logical operations of AND, OR and XOR. Like the arithmetic operations, these can either take two registers as arguments, or one register and an immediate value.

Note that the baseline instruction set does not have a NOT operation to invert all the bits of a register. It's worth considering if that might be useful for your application. One typical way to get around this issue is to XORI with a constant of all 1's, but in our case the 8-bit immediate values are zero-filled in the higher order bits, so that doesn't work. Alternatively you could load up a register with all 1's and then XOR against that register for NOT. Or you could implement a NOT with one of the unused opcodes.

- The Baseline instruction set has one shift instruction: logical shift - a 1-bit left shift or 1-bit right shift. These instructions have an immediate field as part of the instruction. In the full processor design this is used to indicate how many places to shift. A value of 1 in the immediate field shifts the argument register by 1 place to the left. A value of -3 indicates a shift of 3 places to the right. In our Baseline processor, you can assume that all shifts will be by one only. That is, that the only legal values for the immediate are 1 and -1 to indicate left and right shifts. You are, of course, free to augment this in your processor.

An obvious extension of the baseline set would include an arithmetic shift, and that's actually strongly encouraged. Arithmetic shifts are quite useful to perform simple multiplication and division by powers of 2. Arithmetic shifts are the same as Logical shifts except that if shifting is done to the right, the sign bit is copied rather than zero-filled.

- Branch targets are computed as a displacement from the current PC. To compute the branch target a sign-extended offset in the immediate field of the instruction is added to the current

PC and written back to the PC if the branch condition is true. Because the displacement is signed you can branch forwards or backwards from the current PC. But, because the imm field is only 8 bits in our instruction encoding that means that you can only branch to -128 or 127 instructions past the current instruction using the baseline instruction set.

- Jump targets are taken directly from a register. In the case of a jump the register that holds the jump target is written directly to the PC if the jump condition evaluates to true.
- There are two ways to implement Branch, Jump operations. One of them corresponds to using the ALU to compute the target address. The other way would be to design a dedicated ADD/SUB unit for branch target address computations. I would urge you to think about both options and figure out which one would you prefer to design.
- Jump and Link (JAL) is just like a jump, but the PC+1 value is also written to a register. This is known as the link register and lets you jump to a subroutine and then return back to the point in the code where the subroutine was called. You can do this with a JUC Rlink instruction (jump unconditional to the value that you stored in the Rlink register). This is traditionally an instruction that gives people trouble when they implement it. It's not hard, but you do need to get the details right and make sure that eventually you provide a path for the PC to be stored into the register file. That's actually for a later checkpoint though.
- Load and Store instructions need to use values in registers (possibly with immediate values added) as memory addresses. So, you might want a separate path to the Memory Address Register (MAR), or you might put these values through the ALU in some sort of "pass through" mode. Somehow these register values must be usable as memory addresses. Perhaps you can postpone this decision to a future checkpoint too when you implement the MAR and MDR.

Looking at these constraints it is possible to do a general layout for the execution datapath of the machine. The important thing at this point is not to get everything exactly right the first time, but to make sure that you haven't made anything impossible that you will need to do later. The set of control points in this datapath (function codes, mux select signals, latch enables, etc) will be controlled by the control state machine and the decoded instruction.

Register File

The register file must have two read ports. You must be able to read two arguments from the register file to feed the ALU. It also must have a single write port so that you can write back the result. However, since one of the source registers is overwritten with the result, it need only have two addresses specified. These addresses come from the decoded instruction word. You probably want to use Verilog similar to the register file description in the mipscpu.v code for this, but remember that you don't need an independent write address. One of the read addresses is also the write address. You can also look at the bram.v file on the Canvas page for an example of "true dual-port memory" specified as Verilog code.

The ALU should be able to read these two registers' contents and perform an operation on them. However, the ALU must also be able to operate on a register and an immediate value (an ADDI for example), and the PC might also need to be combined with an immediate value (for a displacement-branch for example). These immediates will need to be sign-extended before going into the ALU. You should also consider how to describe sign extension in Verilog (hint - think about concatenation). So muxes on the inputs of the ALU seem required.

Also, it's possible that a register value might be sent to the shifter instead of the ALU. Of course, your ALU might just include a shifting operation and be done with it. Some designs (or designers!) model shifters outside the ALU. In today's world of high-level abstractions and synthesis, such excessive partitioning may be unnecessary. But, if you insist on modeling the shifters outside your ALU, then it is possible that a register value might be sent to the shifter instead of the ALU. So, the output of the execution unit might come from either the ALU or the shifter. Sounds like we may need additional MUXes on inputs/outputs in that case. That's what's shown in Figure 1, but again, that might not be necessary.

One example of a possible organization is shown in Figure 1. *Don't take this block as the required organization!* This is just an example of how things might be organized. It's not even guaranteed to be correct for all the instructions that the Baseline instruction set needs. If you have a different organization in mind, by all means use that. Just make sure that it is able to compute all the values that will be required by the given instruction set! By being clever you may be able to reduce the number of MUXes, or combine units together, or modify things in any number of interesting ways. Another source of inspiration for your RF/ALU organization is the mini-MIPS datapath. Your circuit will be different, of course, because of the difference in instruction sets, but that will be a good source of ideas. This is Figure 1.53 in the Weste/Harris mini-MIPS handout on the Canvas page.

One approach to designing this ALU datapath is to look at every single instruction that you're interested in (i.e. the Baseline instruction set, plus any additional instructions you're thinking about), and check that each one can be accommodated by your ALU datapath.

Some things not shown in Figure 1 are the control points in the datapath. The ALU, for example, will have a function code coming from the control path that tells it which operation it is actually doing at the moment. The MUXes all have select inputs telling them which way they are switched. The latches have enable signals, etc. All these signals are the control points into the execute unit that allows the state machine controller to make it do the right things for each individual instruction. This is all very similar to the `mipscpu.v` that we looked at in the mini-MIPS lab, but of course it's all different because this is a different instruction set!

Remember that your ALU must generate condition codes that will be latched into the processor status register (PSR). Your ALU should generate all the Baseline condition codes. Also pay attention to which instructions update the condition code register! Not all ALU instructions update the PSR register, and even those that do, don't all update it in the same way! Some instructions only set some of the bits, and some instructions don't set any of the bits. Not even all the ADD instructions set the PSR. So, you need to have control over whether the PSR is updated or not. The decoder you design later can decide which PSR bits are updated, but you at least need to plan for the possibility of different sets of bits being set at different times. It's likely that your PSR will be a separate 8-bit register whose outputs can be looked at by subsequent instructions.

What to Do

For this checkpoint you should design an ALU and a register file which will be used eventually in the execute unit of your processor. The overall block will look similar to Figure 1, but definitely will not be exactly like it. Your datapath should support all the instructions from the Baseline instruction set, although you do not have to design the control for those instructions for this lab. You should also be thinking at this point about what types of things you'd like to add to the instruction set. If you already have some ideas, adding support in the data path now might save redesign later. The register file should contain 16 16-bit registers, and your ALU should be capable of performing 16-bit arithmetic and updating the PSR correctly. The program counter and immediate register for this lab will also be 16-bit registers. These will likely be implemented

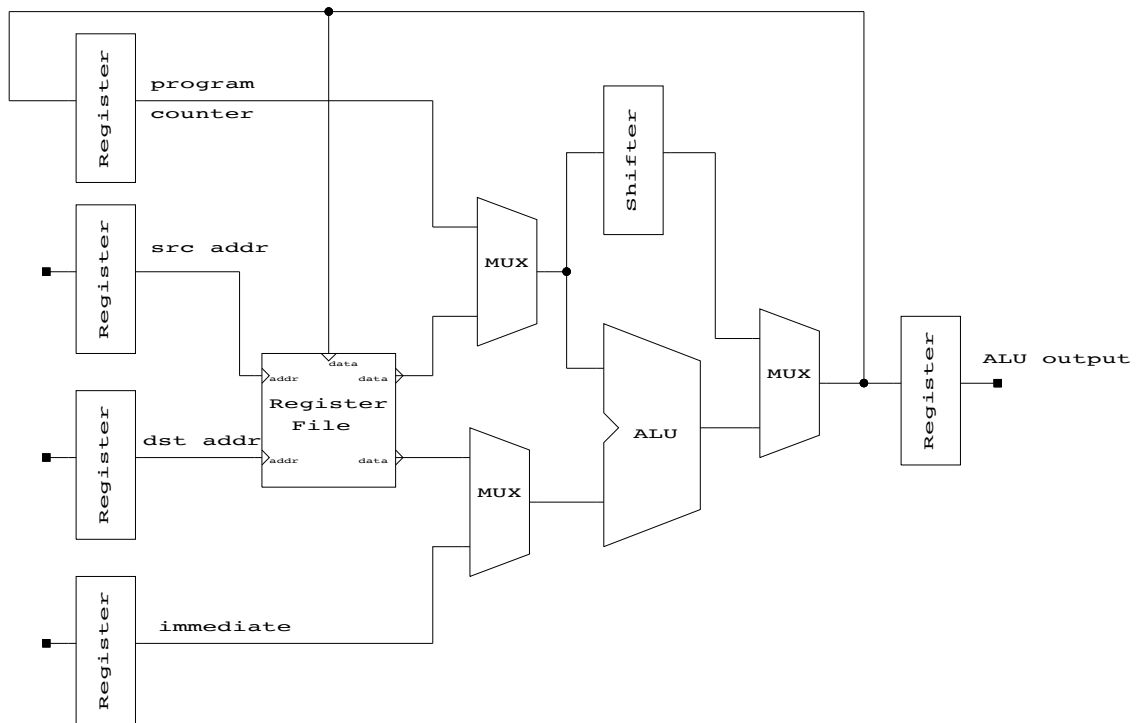


Figure 1: One Possible ALU Organization - just a serving suggestion - your ALU will almost certainly be different!

slightly differently in the final circuit, so make sure that you can change things later.

When you're writing Verilog remember that you should think about the hardware first, and then write Verilog that describes that hardware. Think structurally in terms of system decomposition. Try to partition your design into simple computational units.

Testing

You should test your register file and ALU using the simulator and be prepared to explain your testing procedure, demonstrate the functionality of the circuits, and show the command files and simulation logs from your own testing. Of course, if you can build in checks in the testbench it's much easier to see if things are still working when you make changes! If you have waveforms, make sure you annotate them so we can tell what's going on.

Once you have things simulating, you should synthesize, generate the FPGA file, and upload the circuit to your board. This will actually represent a bit of a problem because without the rest of the processor, it will be tricky to control the RF and ALU. You can use the \$readmemb trick that we used in the mini-MIPS to pre-load the RF with some values, and then perhaps use some of the switches as addresses into the RF. Some of the other switches can be used as ALU function codes. You won't be able to test everything, but you should be able to demonstrate some of the functions of your ALU on at least some pre-set values from the RF.

At least for simulations, the test bench should be as exhaustive as possible to get high coverage and should have self checking capability. The same test bench allows you to verify the circuit and then validate the circuit after you make changes. A good test bench file will also be well commented.

Deliverables

Make a list of all the Opcodes you are implementing, what is the instruction encoding, what addressing modes have you covered so far, etc. Also make a list of all the control points you have identified so far in your RF/ALU circuit. A control point is a MUX select, or a register enable, or an ALU function code, or any other signal that controls some aspect of your RF/ALU so far that will need to be controlled from your eventual finite state control circuit. Generate the above information, collate it into a table - a list of what you implemented will be extremely useful later. You may also want to keep a list of what possible extensions might be needed later.

Upload this document to Canvas along with your Verilog code for the ALU and RF, and test-benches. You should demo your ALU/RF on your DE1-SoC board to the TA or instructor.